

Rationale for MASSExtra

Bill Venables

2020-12-06

Preamble

This extension package to the classical MASS package (Venables & Ripley, of ancient lineage), whose origins go back to nearly 30 years, comes about for a number of reasons.

Firstly, in my teaching I found I was using some of the old functions in the package with consistently different argument settings to the defaults. I was also interested in supplying various convenience extensions that simplified teaching and including various tweaks to improve the interface. Examples follow below.

Secondly, I wanted to provide a few functions that were mainly useful as programming examples. For example, the function `zs` and its allies `zu`, `zq` and `zr` are mainly alternatives to `base::scale`, but they can be used to show how to write functions that can be used in fitting models in such a way that they work as they should when the fitted model object is used for prediction with new data.

Masking select from other packages

Finally, there is the perennial select problem. When MASS is used with other packages, such as `dplyr` the `select` function can easily be masked, causing confusion with users. `MASS::select` is rarely used, but `dplyr::select` is fundamental. There are standard ways of managing this kind of masking, but what we have done in MASSExtra is to export the more common functions used from MASS along with the extensions, in such a way that users will not need to have MASS attached to the search path at all, and hence masking is unlikely.

The remainder of this document will do a walk-through of some of the new functions provided by the package. We begin by setting the computational context:

```
suppressPackageStartupMessages({
  library(ggwebthemes) ## https://gitlab.com/peterbar/ggwebthemes/
  library(visreg)
  library(knitr)
  library(tidyverse)
  library(patchwork)
  library(MASSExtra)
})
options(knitr.kable.NA = "")
theme_set(theme_web_bw() + theme(title = element_text(hjust = 0.5)))
```

Amble

We now consider some of the extensions that the package offers to the originals. Most of the extensions will have a name that includes an underscore of two somewhere to distinguish it from the V&R original. Note that the original version is *also* exported so that scripts that use it may do so without change, via the new package.

The `box_cox` extensions

This original version, `boxcox` has a fairly rigid display for the plotted output which has been changed to give a more easily appreciated result. The y -axis has been changed to give the likelihood-ratio statistic rather

than the log-likelihood, and for the x -axis some attempt has been made to focus on the crucial region for the transformation parameter, λ ,

The following example shows the old and new plot versions for a simple example.

```
par(mfrow = c(1, 2))
mod0 <- lm(MPG.city ~ Weight, Cars93)
boxcox(mod0) ## MASS
box_cox(mod0) ## MASSExtra tweak
```




Figure 1: Box-cox, old and new displays

In addition, there are functions `bc` to evaluate the transformation for a given exponent, and a function `lambda` which finds the optimum exponent (not that a precise exponent will usually be needed).

It is interesting to see how in this instance the transformation can both straighten the relationship and provide a scale in which the variance is more homogeneous. See Figure 2.

```
p0 <- ggplot(Cars93) + aes(x = Weight) + geom_point(colour = "#2297E6") + xlab("Weight (lbs)") +
  geom_smooth(se = FALSE, method = "loess", formula = y ~ x, size=0.7, colour = "black")
p1 <- p0 + aes(y = MPG.city) + ylab("Miles per gallon (MPG)") + ggtitle("Untransformed response")
p2 <- p0 + aes(y = bc(MPG.city, lambda(mod0))) + ggtitle("Transformed response") +
  ylab(bquote(bc(MPG, .(round(lambda(mod0), 2)))))
p1 + p2
```

A more natural scale to use, consistent with the Box-Cox suggestion, would be the reciprocal. For example we could use $GPM = 100/MPG$ the “gallons per 100 miles” scale, which would have the added benefit of being more-or-less what the rest of the world uses to gauge fuel efficiency outside the USA. Readers should try this for themselves.

Stepwise model building extensions

The primary MASS functions for refining linear models and their allies are `dropterm` and `stepAIC`. The package provides a few extensions to these, but mainly a change of defaults in the argument settings.

1. `drop_term` is a front-end to `MASS::dropterm` with a few tweaks. By default the result is arranged in sorted order, i.e. with `sorted = TRUE`, and also by default with `test = TRUE` (somewhat in defiance of much advice to the contrary given by experienced practitioners: *caveat emptor!*).

The user may specify the test to use in the normal way, but the default test is decided by an ancillary

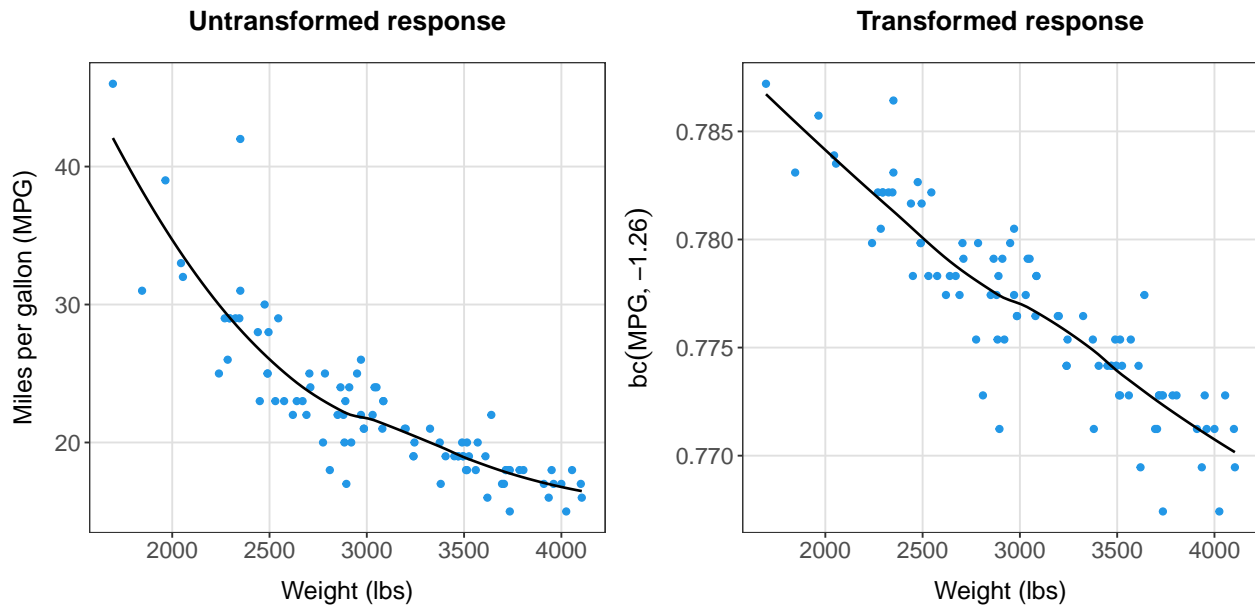


Figure 2: The Box-Cox transformation effect

generic function, `default_test`, which guesses the appropriate test from the object itself. This is an S3 generic and further methods can be supplied for new fitted model objects.

In addition `drop_term` returns an object which retains information on the criterion used, AIC, BIC, GIC (see below) or some specific penalty value k . The object also has a class "drop_term" for which a plot method is provided. Both the plot and print methods display the criterion. See the example below for how this is done.

2. `step_AIC` is a front-end to `MASS::stepAIC` with the default argument `trace = FALSE` set. This may of course be over-ruled, but it seems the most frequent choice by users, anyway. In addition the actual criterion used, by default $k = 2$, i.e. AIC, is retained with the result and passed on to methods in much the same way as for `drop_term` above.

Since the (default) criterion name is encoded in the function name, two further versions are supplied, namely `step_BIC` and `step_GIC` (again, see below), which use a different, and obvious, default criterion.

In any of `step_AIC`, `step_BIC` or `step_GIC` a different value of k may be specified in which case that value of k is retained with the object and displayed as appropriate in further methods.

Finally in any of these functions k may be specified either as a numeric penalty, such as $k = 4$ for example, or by character string $k = \text{"AIC"}$ or $k = \text{"BIC"}$ with an obvious meaning in either case.

3. **Criteria.** The **Akaike Information Criterion**, AIC, corresponds to a penalty $k = 2$ and the **Bayesian Information Criterion**, BIC, corresponds to $k = \log(n)$ where n is the sample size. In addition to these two the present functions offer an intermediate default penalty $k = (2 + \log(n))/2$ which is "not too strong and not too weak", making it the **Goldilocks Information Criterion**, GIC. There is also a standalone function `GIC` to evaluate this k if need be.

This suggestion appears to be original, but *no particular claim is made for it* other than with intermediate to largish data sets it has proved useful for exploratory purposes in our experience.

Our strong advice is that these tools should *only* be used for exploratory purposes in any case, and should *never* be used in isolation. They have a well-deserved very negative reputation when misused, as they commonly are.

Examples

We consider the well-known and much maligned Boston house price data.